


Object Recognition with CNNs & Transfer Learning (CIFAR-10)



Class: airplane



Class: ship



Class: cat



Class: frog



Class: deer



Class: frog



Class: cat



Class: airplane



Class: automobile



Class: dog



Class: horse




Class: cat



Class: truck




Class: automobile



Class: ship



Class: bird



Dataset & Pre- processing

CIFAR-10: 60k images, 10
classes, $32 \times 32 \times 3$

Train/Val/Test split:
45k/5k/10k

Normalisation, one-hot
encoding, consistent
preprocessing



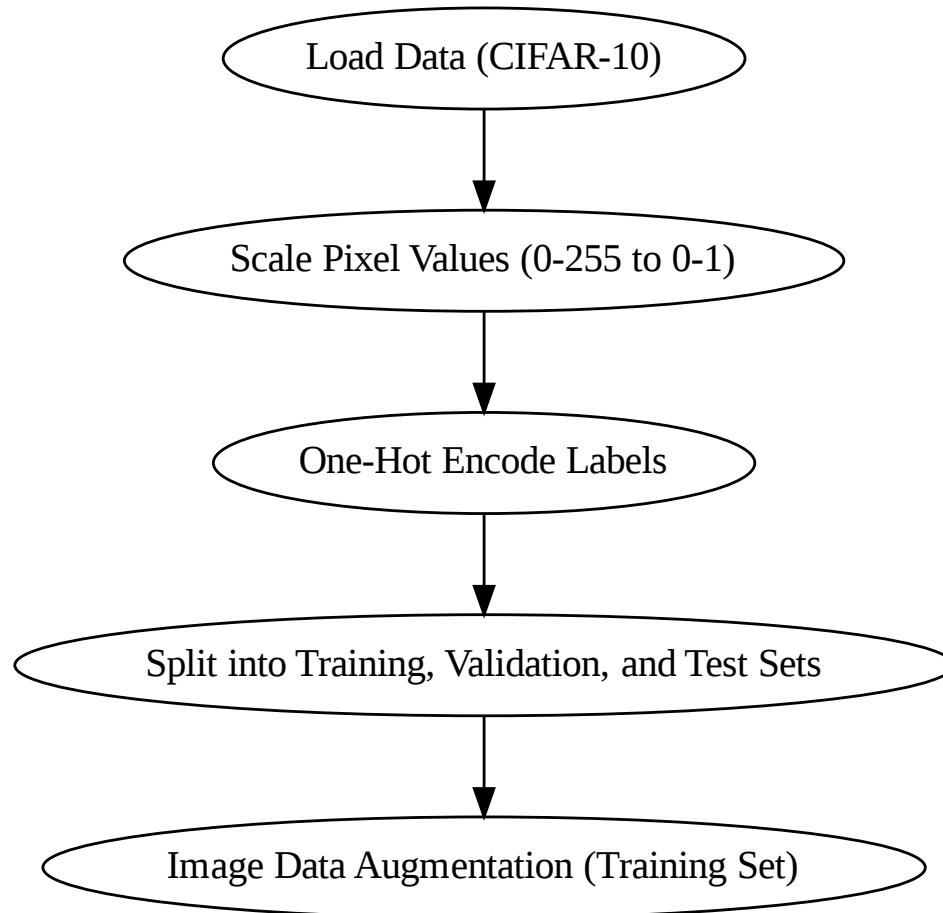
Data Augmentation

Random flip, rotation,
zoom, translation

Enhances diversity
and generalisation

Applied during
training to both
models






Model 1 –
Baseline
CNN (from
scratch)

Conv–Conv–Pool ×2 →
Conv → Dense(256) →
Dropout → Softmax(10)

Optimizer: Adam (1e-3),
Loss: CCE, Dropout
regularisation





--- Custom CNN Model Summary (without data augmentation) ---

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 29, 29, 32)	1,568
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_3 (Conv2D)	(None, 11, 11, 32)	16,416
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 32)	0
flatten_2 (Flatten)	(None, 800)	0
dense_4 (Dense)	(None, 256)	205,056
dense_5 (Dense)	(None, 10)	2,570

Total params: 676,832 (2.58 MB)

Trainable params: 225,610 (881.29 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 451,222 (1.72 MB)



Model 2 – Transfer Learning (VGG16 fine-tune)

Base: VGG16 pretrained on
ImageNet, include_top=False


Freeze early layers; fine-tune
last block

Custom head: GAP →
Dense(256) → Dropout →
Softmax(10)





VGG-16 CNN Architecture



--- VGG16 Model Summary ---

Model: "sequential_3"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 3, 3, 512)	14,714,688
flatten_3 (Flatten)	(None, 4608)	0
dense_6 (Dense)	(None, 256)	1,179,904
dense_7 (Dense)	(None, 10)	2,570

Total params: 18,262,112 (69.66 MB)

Trainable params: 1,182,474 (4.51 MB)

Non-trainable params: 14,714,688 (56.13 MB)

Optimizer params: 2,364,950 (9.02 MB)




Training Curves (Accuracy/Loss)

Compare CNN vs VGG16
training & validation
accuracy

VGG16 converges faster;
CNN may overfit earlier



evaluation_accuracy_vs_iterations



26250 x 2

Run ↑

Smoothed

Value

Step

Relative



cnn_aug/validation

0.7065

0.7146

26,250

7.766 min



cnn_no_aug/validation

0.6645

0.6773

12,500

39.85 sec



vgg16/validation

0.7799


0.7856

16,250

17.45 min



evaluation_loss_vs_iterations



Run ↑

Smoothed

Value

Step

Relative

●	cnn_aug/validation	0.864	0.8507	26,250	7.766 min
●	cnn_no_aug/validation	1.0381	1.0369	12,500	39.85 sec
●	vgg16/validation	0.6428	0.6355	16,250	17.45 min



Results & Confusion Matrices

Validation/Test Accuracy
for both models

Confusion matrices
highlight misclassifications

VGG16 reduces class
confusion vs CNN



Model	Data Augmentation	Test Loss	Test Accuracy
Custom CNN	No	1.0458	67.20%
Custom CNN	Yes	1.0458	70.78%
VGG16	Yes	0.6369	78.03%





Commands + Code + Text ▶ Run all

✓ RAM
Disk

```
32 except Exception as e:
33     print(f"Error loading model from {model_aug_path}: {e}")
34     loaded_model_aug = None # Set to None if loading fails
35
36 try:
37     loaded_model_vgg16 = tf.keras.models.load_model(model_vgg16_path)
38     print(f"Model loaded successfully from: {model_vgg16_path}")
39 except Exception as e:
40     print(f"Error loading model from {model_vgg16_path}: {e}")
41     loaded_model_vgg16 = None # Set to None if loading fails
```

Mounted at /content/drive
Model loaded successfully from: /content/drive/MyDrive/saved_models/cnn_no_aug.keras
Model loaded successfully from: /content/drive/MyDrive/saved_models/cnn_aug.keras
Model loaded successfully from: /content/drive/MyDrive/saved_models/vgg16_cifar10.keras

[11]
✓ Os

```
1 # Evaluate Custom CNN without Data Augmentation
2 print("\nEvaluating Custom CNN without Data Augmentation:")
3 # Assuming 'model_cnn' was trained without augmentation in the last run of cell 0ce4775e
4 loss_no_aug, accuracy_aug = loaded_model_no_aug.evaluate(x_test, y_cat_test, verbose=0)
5 print(f"Test Loss: {loss_no_aug:.4f}")
6 print(f"Test Accuracy: {accuracy_aug:.4f}")
```

Evaluating Custom CNN without Data Augmentation:
Test Loss: 1.0458
Test Accuracy: 0.6720

[]

```
1 # Evaluate Custom CNN with Data Augmentation
2 print("\nEvaluating Custom CNN with Data Augmentation:")
3 loss_no_aug, accuracy_aug = loaded_model_aug.evaluate(x_test, y_cat_test, verbose=0)
4 print(f"Test Loss: {loss_no_aug:.4f}")
5 print(f"Test Accuracy: {accuracy_aug:.4f}")
```

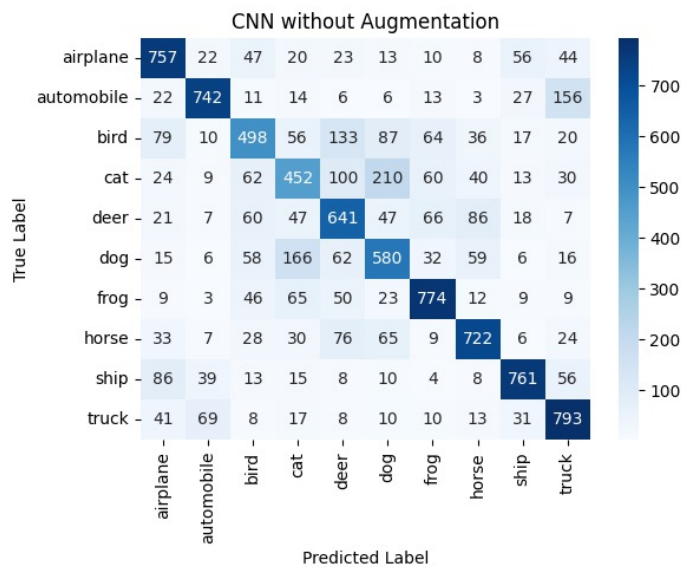
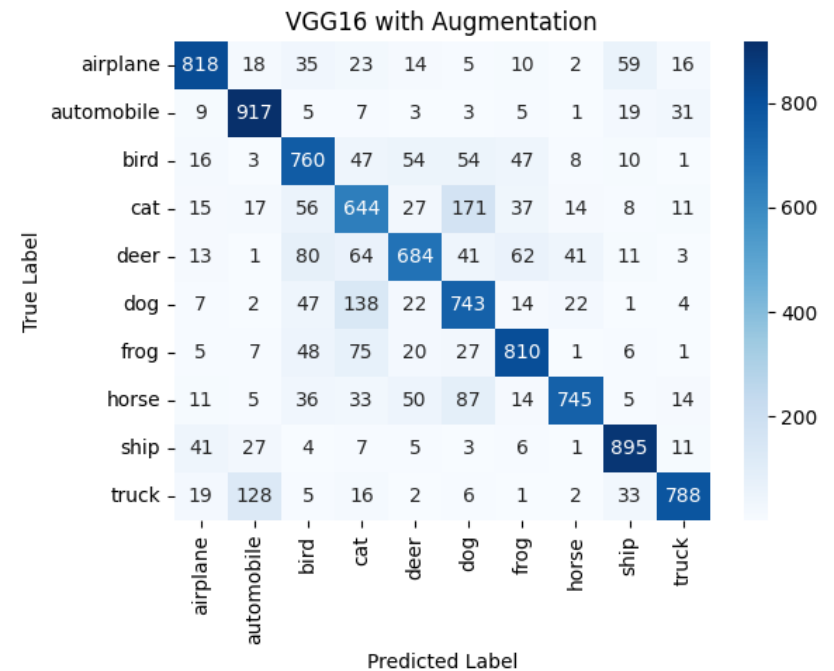
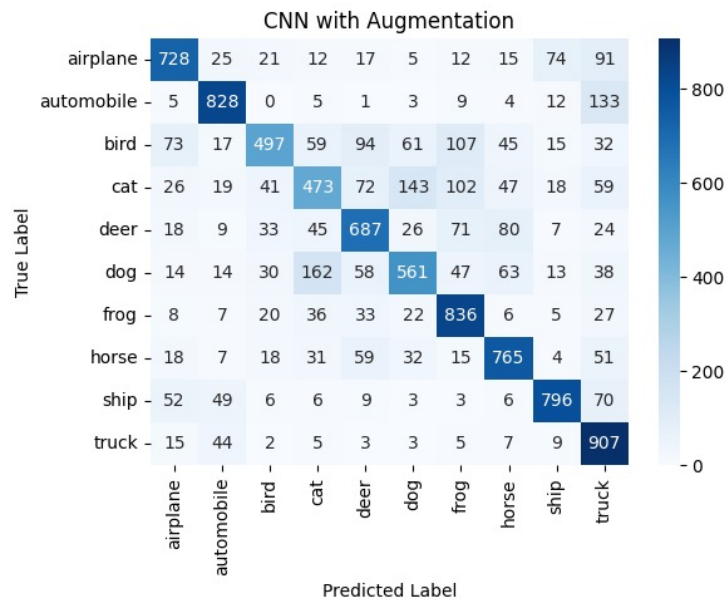
Evaluating Custom CNN with Data Augmentation:
Test Loss: 1.0458
Test Accuracy: 0.7078

[]

```
1 # Evaluate VGG16 with Data Augmentation
2 print("\nEvaluating VGG16 with Data Augmentation:")
3 # Assuming 'model_vgg16' was trained with augmentation in the last run of cell 8dfb978d
4 loss_vgg16, accuracy_vgg16 = loaded_model_vgg16.evaluate(x_test_resized, y_cat_test_vgg16, verbose=0)
5 print(f"Test Loss: {loss_vgg16:.4f}")
6 print(f"Test Accuracy: {accuracy_vgg16:.4f}")
```

Evaluating VGG16 with Data Augmentation:
Test Loss: 0.6369
Test Accuracy: 0.7803

Double-click (or enter) to edit



Class	Precision	Recall	F1-Score	Support
airplane	0.7	0.76	0.73	1000
automobile	0.81	0.74	0.78	1000
bird	0.6	0.5	0.54	1000
cat	0.51	0.45	0.48	1000
deer	0.58	0.64	0.61	1000
dog	0.55	0.58	0.57	1000
frog	0.74	0.77	0.76	1000
horse	0.73	0.72	0.73	1000
ship	0.81	0.76	0.78	1000
truck	0.69	0.79	0.74	1000
Accuracy	nan	nan	0.67	10000
Macro avg	0.67	0.67	0.67	10000
Weighted avg	0.67	0.67	0.67	10000

CNN no augmentation

Class	Precision	Recall	F1-Score	Support
airplane	0.76	0.73	0.74	1000
automobile	0.81	0.83	0.82	1000
bird	0.74	0.5	0.6	1000
cat	0.57	0.47	0.52	1000
deer	0.67	0.69	0.68	1000
dog	0.65	0.56	0.6	1000
frog	0.69	0.84	0.76	1000
horse	0.74	0.77	0.75	1000
ship	0.84	0.8	0.82	1000
truck	0.63	0.91	0.75	1000
Accuracy	nan	nan	0.71	10000
Macro avg	0.71	0.71	0.7	10000
Weighted avg	0.71	0.71	0.7	10000

CNN with augmentation

Class	Precision	Recall	F1-Score	Support
airplane	0.86	0.82	0.84	1000
automobile	0.82	0.92	0.86	1000
bird	0.71	0.76	0.73	1000
cat	0.61	0.64	0.63	1000
deer	0.78	0.68	0.73	1000
dog	0.65	0.74	0.69	1000
frog	0.81	0.81	0.81	1000
horse	0.89	0.74	0.81	1000
ship	0.85	0.9	0.87	1000
truck	0.9	0.79	0.84	1000
Accuracy	nan	nan	0.78	10000
Macro avg	0.79	0.78	0.78	10000
Weighted avg	0.79	0.78	0.78	10000

VGG16



Discussion – Strengths & Trade-offs



CNN: lightweight,
flexible, lower
compute



Transfer Learning:
higher accuracy,
faster convergence,
heavier mode



Lessons Learned & Insights



Hyperparameter
tuning



Importance of
Validation Set



Power of data-
augmentation



Applications, Ethics & Future Work



Autonomous Vehicles,
Biometric Authentication,
Medical Imaging




Ethical and practical
dimensions of AI



ResNet50/MobileNetV2





<https://stanfordmag.org/contents/in-two-years-there-could-be-10-million-self-driving-cars-on-the-roads>



References

- Bender, E.M., Gebru, T., McMillan-Major, A. and Shmitchell, S. (2021) *On the dangers of stochastic parrots: Can language models be too big?* In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FACt '21)*. New York: ACM Press.
- Esteva, A., et al. (2019) *A guide to deep learning in healthcare*. *Nature Medicine*, 25, pp.24–29.
- Goodfellow, I., Bengio, Y. and Courville, A. (2016) *Deep learning*. Cambridge, MA: MIT Press.
- Hien, B. M. (2023) *Solving CIFAR-10 dataset with VGG16 pre-trained architect using PyTorch*. Available at: <https://medium.com/@buiminhkien2k/solving-cifar10-dataset-with-vgg16-pre-trained-architect-using-pytorch-validation-accuracy-over-3f9596942861>.
- Howard, J. and Ruder, S. (2018) *Universal language model fine-tuning for text classification*. In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)*. Melbourne: ACL, pp.328–339.
- Jobin, A., Ienca, M. and Vayena, E. (2019) *The global landscape of AI ethics guidelines*. *Nature Machine Intelligence*, 1(9), pp.389–399.
- Kingma, D.P. and Ba, J. (2015) *Adam: A method for stochastic optimization*. In: *Proceedings of the International Conference on Learning Representations (ICLR 2015)*. Available at: <https://arxiv.org/abs/1412.6980>
- Krizhevsky, A., Nair, V. and Hinton, G. (2009) *The CIFAR-10 dataset*. Toronto: University of Toronto. Available at: <https://www.cs.toronto.edu/~kriz/cifar.html>



References Continued

LeCun, Y., Bengio, Y. and Hinton, G. (2015) *Deep learning*. *Nature*, 521(7553), pp.436–444.

LeCun, Y., Bottou, L., Bengio, Y. and Haffner, P. (1998) *Gradient-based learning applied to document recognition*. *Proceedings of the IEEE*, 86(11), pp.2278–2324.

Shorten, C. and Khoshgoftaar, T.M. (2019) *A survey on image data augmentation for deep learning*. *Journal of Big Data*, 6(1), 60.

Simonyan, K. and Zisserman, A. (2015) *Very deep convolutional networks for large-scale image recognition*. *International Conference on Learning Representations (ICLR 2015)*. Available at: <https://arxiv.org/abs/1409.1556>

Yosinski, J., Clune, J., Bengio, Y. and Lipson, H. (2014) *How transferable are features in deep neural networks?* In: *Proceedings of the 27th International Conference on Neural Information Processing Systems (NeurIPS 2014)*. Montréal, Canada.

Zhang, Y., Bengio, S., Hardt, M., Recht, B. and Vinyals, O. (2021) *Understanding deep learning requires rethinking generalization*. *International Conference on Learning Representations (ICLR 2021)*. Available at: <https://arxiv.org/abs/1611.03530>

